

CSE4261: Neural Network and Deep Learning

Lecture: 24.06.2025



Sangeeta Biswas, Ph.D.
Associate Professor,
University of Rajshahi, Rajshahi-6205, Bangladesh

Gradient Descent Algorithm for Parameter Updating

Gradient is calculated for cost function with respect to weights and biases for updating their values.

$$w^{[l]} = w^{[l]} - \alpha \frac{\partial C}{\partial w^{[l]}}$$

$$b^{[l]} = b^{[l]} - \alpha \frac{\partial C}{\partial b^{[l]}}$$

Adversarial Attack by Fast Gradient Signed Method (FGSM)

- FGSM uses the gradients of the loss with respect to the input image to create a new image that maximises the loss.
- Adversarial image, $adv_x = x + \epsilon * \text{sign}(\nabla_x J(\theta, x, y))$




x

“panda”

57.7% confidence

+ .007 ×



$\text{sign}(\nabla_x J(\theta, x, y))$

“nematode”

8.2% confidence

=

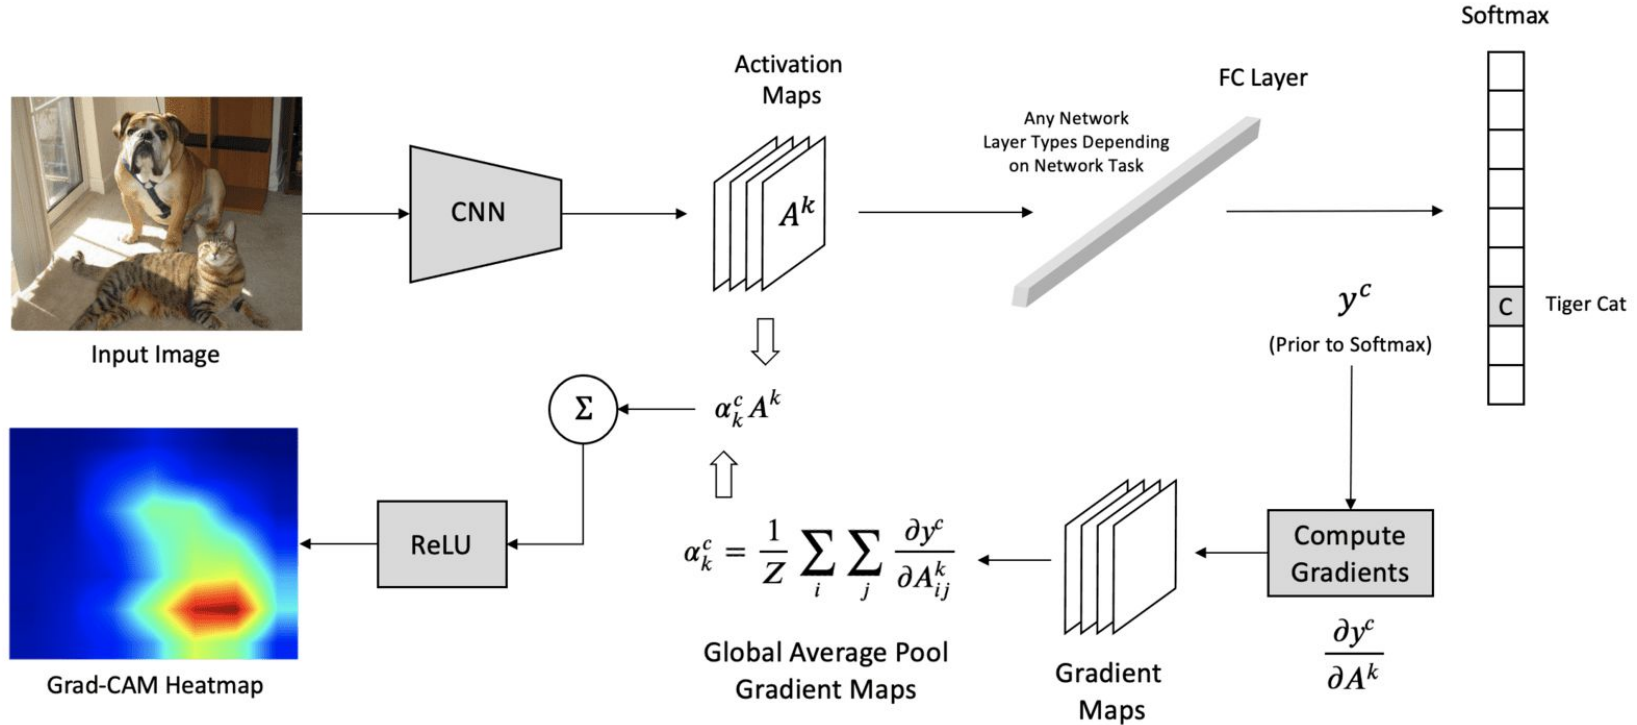


$x + \epsilon \text{sign}(\nabla_x J(\theta, x, y))$

“gibbon”


99.3 % confidence

Grad-CAM (Gradient weighted Class Activation Mapping)




Interpolation

- It is a method to estimate the value of a function at a point within the range of a known set of data points.
- Given a set of data points $(x_1, y_1), \dots, (x_n, y_n)$, interpolation finds the value of y for a given x that lies within the range of x_1 and x_n




Linear Interpolation

- It finds an intermediate value between two known values by assuming a linear relationship between them.




Integrated Gradients (IG) Method as an XAI Technique

- In IG, a baseline image is used with the input image.
- Linear interpolation is used to approximate the integral of gradients along a path between a baseline and the input instance (i.e., path P_2).
- The integral is approximated by summing the gradients at multiple interpolated points along this line.



Integrated Gradients (IG) Method as an XAI Technique


Baseline Image



Original Image



IG Attribution Mask



Original + IG Attribution Mask Overlay




- Tutorial:
https://www.tensorflow.org/tutorials/interpretability/integrated_gradients

Example of Linear Interpolation

- Five-step interpolation between the baseline x' and the input image x

$\alpha = 0.0$



Baseline

$\alpha = 0.2$



$\alpha = 0.4$



$\alpha = 0.6$



$\alpha = 0.8$



$\alpha = 1.0$



Input image

Formula of IG

$$\text{IntegratedGrads}_i^{\text{approx}}(x) ::= (x_i - x'_i) \times \sum_{k=1}^m \frac{\partial F(x' + \frac{k}{m} \times (x - x'))}{\partial x_i} \times \frac{1}{m}$$

i = feature (individual pixel)

x = input (image tensor)

x' = baseline (image tensor)

k = scaled feature perturbation constant

m = number of steps in the Riemann sum approximation of the integral

(x_i-x'_i)= a term for the difference from the baseline

F = model prediction function

Formula of IG

$$\textit{IntegratedGrads}_i^{\textit{approx}}(x) ::= (x_i - x'_i) \times \sum_{k=1}^m \overbrace{\frac{\partial F(\textit{interpolated images})}{\partial x_i}}^{\text{compute gradients}} \times \frac{1}{m}$$

$$\frac{\partial F}{\partial x_i} = \text{gradient}$$

- The gradient tells us which pixels have the strongest effect on the model's predicted class probabilities.

Gradient Computation for IG by Tensorflow

```
def compute_gradients(images, target_class_idx):  
    with tf.GradientTape() as tape:  
        tape.watch(images)  
        logits = model(images)  
        probs = tf.nn.softmax(logits, axis = -1)[:, target_class_idx]  
    return tape.gradient(probs, images)
```

In Grad-CAM:

```
model.layers[-1].activation = None # Remove last layer's softmax
```

Alternative Baseline

Baseline - Add Gaussian



(a) Gaussian Baseline

Baseline - Blur



(b) Blur Baseline

Baseline - Max Distance



(c) Max Distance Baseline

Baseline - Uniform



(d) Uniform Baseline